

1703.03400 - Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks

- **Yunqiu Xu**
- 2nd for this paper, try to understand both paper and code
- Other reference:
 - <https://www.jiqizhixin.com/articles/2017-07-20-4>
 - <http://bair.berkeley.edu/blog/2017/07/18/learning-to-learn/>
- Implementation:
 - Original: <https://github.com/cbfinn/maml>
 - PyTorch: <https://github.com/katerakelly/pytorch-maml>

1. Introduction

- The goal of meta-learning: train a model with some learning tasks, then it can solve new tasks with only a few samples
- Our work:
 - MAML: pretrain the model with a series of tasks
 - Then this model can be generated to new task with a small number of training samples / gradient steps on that task
- Comparison with another recent related work: **Ravi & Larochelle, 2017.**

Optimization as a Model for Few-Shot Learning

- Their work :
 - Learns both the weight initialization and the optimizer → they learn an update function or learning rule
 - Meta-based LSTM
- Our advantage :
 - The MAML learner' s weights are updated using the gradient, not a trained update method
 - So we do not need additional parameters
 - And we are not limited to LSTM, our model can be generalized to both SL

and RL

- **Other advantages:** see 6.Discussion and Future Work

2. MAML

2.1 Problem Setup

- Goal: train a "meta learning" model on a set of tasks, then this model can adapt to a new task with only a few data / iterations → learn **as much as** possible with limited data
- Model $f: x \rightarrow a$
 - x : observations
 - a : outputs
 - This model is like a "base model" which will be able to adapt to a lot of new tasks
- General notion of task $T = \{L(x_1, a_1, \dots, x_H, a_H), q(x_1), q(x_{t+1}|x_t, a_t), H\}$:
 - $L \rightarrow R$: loss function
 - $q(x_1)$: distribution over initial observations
 - $q(x_{t+1}|x_t, a_t)$: transition distribution
 - 对于监督学习, 不存在这个分布: $H = 1$
 - 对于强化学习, $q(x_{t+1}|x_t, a_t)$ 代表某时间点观察值的分布, e.g. 初始观察值后观察值取自 $q(x_2|x_1, a_1)$
 - H :
 - Episode length, model may generate samples of length H by choosing an output a_t at each time t
 - For supervised learning, $H = 1$ and loss function $L(x_1, a_1)$ could be MSE or cross entropy

2.2 K-shot meta-learning:

- Train model f :
 - Sample a new task T_i from $p(T)$ (training taskset)

- Learn T_i :
 - Train model with K samples drawn from q_i
 - Get feedback L_{T_i} from T_i
- Test on new samples from T_i and get test error
- Improve model f : treat the test error on sampled tasks T_i as the training error of meta-learning process
- Test meta-learning:
 - Sample new task from $p(T)$ (testing taskset), try to adapt f to this new task
 - Learn the model with K samples
 - Treat the performance as "meta-performance"

2.3 A MAML Algorithm

- What our model does?
 - Be able to learn parameters of any standard model via meta-learning
 - Why: some internal representations are more transferrable
 - E.G. 我们可以通过一系列任务学到一个神经网络模型, 而非仅仅通过一个任务, 这样这个模型就比较容易迁移到类似的新任务上
- How we learn
 - Learn a model that gradient-based learning rule can make rapid progress on new tasks drawn from $p(T)$ without overfitting
 - Find model parameters that are **sensitive** to changes of the task → small changes lead to large improvement (direction of gradient)
 - 为什么这样做: 之前的目标就是在样本/更新次数有限的情况下学到的东西尽可能多, 因此我们要尽量让每一点点小改变都能获得较大的提升
 - **此处存疑: 代码里该怎么体现"sensitive", 看了代码好像没有具体提及**




Figure 1. Diagram of our model-agnostic meta-learning algorithm (MAML), which optimizes for a representation θ that can quickly adapt to new tasks.

• Algorithm 1:

- f_θ , when adapting to a new training task T_i , the model's parameter θ becomes θ'_i
- Then we learn θ'_i by gradient update:

$$\theta'_i = \theta - \alpha \nabla_{\theta} L_{T_i}(f_\theta)$$

- 这里我理解为调整梯度的方向, 比如为了适应 T_2 我们需要将梯度方向稍微上移
- 经过多次梯度更新(内循环)后, 我们就可以学到比较适合 T_2 的参数向量 θ'_2
- Step size α : fixed as a hyperparameter or meta-learned
- After learning the parameter vectors for all tasks in training task set ($\theta'_1, \theta'_2, \dots, \theta'_n$), compute their test error to update θ :

- Meta objective

$$\min_{\theta} \sum_{T_i \sim p(T)} L_{T_i}(f_{\theta'_i}) = \sum_{T_i \sim p(T)} L_{T_i}(f_{\theta - \alpha \nabla_{\theta} L_{T_i}(f_{\theta})})$$

- SGD:

$$\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{T_i \sim p(T)} L_{T_i}(f_{\theta'_i})$$

Algorithm 1 Model-Agnostic Meta-Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

- 1: randomly initialize θ
 - 2: **while** not done **do**
 - 3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$
 - 4: **for all** \mathcal{T}_i **do**
 - 5: Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ with respect to K examples
 - 6: Compute adapted parameters with gradient descent: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$
 - 7: **end for**
 - 8: Update $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$
 - 9: **end while**
-

- This process involves a gradient through a gradient \rightarrow an additional backward pass through f to compute Hessian-vector products

3. Species of MAML

Algorithm 2 MAML for Few-Shot Supervised Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

- 1: randomly initialize θ
 - 2: **while** not done **do**
 - 3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$
 - 4: **for all** \mathcal{T}_i **do**
 - 5: Sample K datapoints $\mathcal{D} = \{\mathbf{x}^{(j)}, \mathbf{y}^{(j)}\}$ from \mathcal{T}_i
 - 6: Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ using \mathcal{D} and $\mathcal{L}_{\mathcal{T}_i}$ in Equation (2) or (3)
 - 7: Compute adapted parameters with gradient descent: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$
 - 8: Sample datapoints $\mathcal{D}'_i = \{\mathbf{x}^{(j)}, \mathbf{y}^{(j)}\}$ from \mathcal{T}_i for the meta-update
 - 9: **end for**
 - 10: Update $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$ using each \mathcal{D}'_i and $\mathcal{L}_{\mathcal{T}_i}$ in Equation 2 or 3
 - 11: **end while**
-

Algorithm 3 MAML for Reinforcement Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

- 1: randomly initialize θ
 - 2: **while** not done **do**
 - 3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$
 - 4: **for all** \mathcal{T}_i **do**
 - 5: Sample K trajectories $\mathcal{D} = \{(\mathbf{x}_1, \mathbf{a}_1, \dots, \mathbf{x}_H)\}$ using f_{θ} in \mathcal{T}_i
 - 6: Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ using \mathcal{D} and $\mathcal{L}_{\mathcal{T}_i}$ in Equation 4
 - 7: Compute adapted parameters with gradient descent: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$
 - 8: Sample trajectories $\mathcal{D}'_i = \{(\mathbf{x}_1, \mathbf{a}_1, \dots, \mathbf{x}_H)\}$ using $f_{\theta'_i}$ in \mathcal{T}_i
 - 9: **end for**
 - 10: Update $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$ using each \mathcal{D}'_i and $\mathcal{L}_{\mathcal{T}_i}$ in Equation 4
 - 11: **end while**
-

3.1 Supervised learning

- Model is to predict output value given input value
- For each training task T_i
 - $H = 1$: single input and single output
 - $L = \text{MSE or corss entropy}$
 - $q_i(x_1)$: 因为不存在时序观察值, 这个分布就是监督学习训练集样本的分布
 - I think there is no $q_i(x_{t+1}|x_t, a_t)$
 - Then generate K samples x from $q_i(x_1)$, compute the error between predicted value a and ground truth y

3.2 RL

- Model is to predict action a_t given state x_t
- For each training task T_i :
 - timestep $t \in \{1, \dots, H\}$
 - The initial state distribution $q_i(x_1)$: 为了训练这个子任务, 我们会从这个初始值分布选取K个起始点
 - Transition distribution $q_i(x_{t+1}|x_t, a_t)$:
 - 对每个当前观察值及选取的动作, 未来观察值同样构成一个分布
 - 比如我选择吃一口饭, 接下来可能观察到 {饱了, 还饿} 等状态
 - 当然我们训练好 T_i , 获得的未来状态可能就是选取最优动作后的结果了
 - For T_i and its model parameter ϕ , the loss function is

$$L_{T_i}(f_\psi) = -E_{x_t, a_t \sim f_{\psi, q_{T_i}}} \left[\sum_{t=1}^H R_i(x_t, a_t) \right] \quad (4)$$

- As reward function is to maximize reward, in loss function we multiply "-1" to minimize the value
- Here ϕ is θ'_i , which we mentioned before
- Why in step 8 we sample trajectories using $f_{\theta'_i}$ instead of f_{θ_i} : PG is on-policy, thus each additional gradient updating during the adaption of f_θ need to sample from current policy

4. Related work

	5-way Accuracy		20-way Accuracy	
	1-shot	5-shot	1-shot	5-shot
Omniglot (Lake et al., 2011)				
MANN, no conv (Santoro et al., 2016)	82.8%	94.9%	–	–
MAML, no conv (ours)	89.7 ± 1.1%	97.5 ± 0.6%	–	–
Siamese nets (Koch, 2015)	97.3%	98.4%	88.2%	97.0%
matching nets (Vinyals et al., 2016)	98.1%	98.9%	93.8%	98.5%
neural statistician (Edwards & Storkey, 2017)	98.1%	99.5%	93.2%	98.1%
memory mod. (Kaiser et al., 2017)	98.4%	99.6%	95.0%	98.6%
MAML (ours)	98.7 ± 0.4%	99.9 ± 0.1%	95.8 ± 0.3%	98.9 ± 0.2%

	5-way Accuracy	
	1-shot	5-shot
MiniImagenet (Ravi & Larochelle, 2017)		
fine-tuning baseline	28.86 ± 0.54%	49.79 ± 0.79%
nearest neighbor baseline	41.08 ± 0.70%	51.04 ± 0.65%
matching nets (Vinyals et al., 2016)	43.56 ± 0.84%	55.31 ± 0.73%
meta-learner LSTM (Ravi & Larochelle, 2017)	43.44 ± 0.77%	60.60 ± 0.71%
MAML, first order approx. (ours)	48.07 ± 1.75%	63.15 ± 0.91%
MAML (ours)	48.70 ± 1.84%	63.11 ± 0.92%

- RNNs as learners: MANN
 - Search space includes all conceivable ML algorithms
 - Moves the burden of innovation to RNNs
 - Ignores advances achieved in ML by humans
 - The results are not good
- Metric learning: Siamese nets, matching nets
 - Learn a metric in input space
 - Specialized to one/few-shot classification
 - Can't use in other problems
- Optimizer learning: meta-learner LSTM
 - Learn parameter update given gradients (search space includes SGD, RMSProp, Adam etc)
 - Applicable to any architecture / task
 - But we can achieve better performance with MAML

5. Experimental evaluation

- Questions need to be answered:

- Can MAML enable fast learning of new tasks
- Can MAML be used for meta-learning in multiple different domains: SL / RL ...
- Can a model learned with MAML continue to improve with additional gradient updates and/or examples
- An oracle work:
 - Receives the identity of the task (which is a problem-dependent representation) as an additional input,
 - Thus oracle will be an upper bound on the performance of the model
 - 我们主要对比用或不用MAML的情形, 越接近oracle代表效果越好

5.1 Regression

- 模拟sin曲线
- Each task:
 - The shape of curve varies by amplitude and phase
 - Input and output of a sine wave
 - Data points sampled from $[-5.0, 5.0]$
 - Loss: MSE
- Model architecture: NN with 2 hidden layers, 40 hidden nodes, ReLU
- Training:
 - $K = 10$
 - $\alpha = 0.01$
 - Adam
 - After all training tasks, we get a pretrained model
- Testing:
 - Fine-tune the pretrained model with K test samples and a number of GDs
- Result
 - Left is pretrained with MAML, right is without MAML

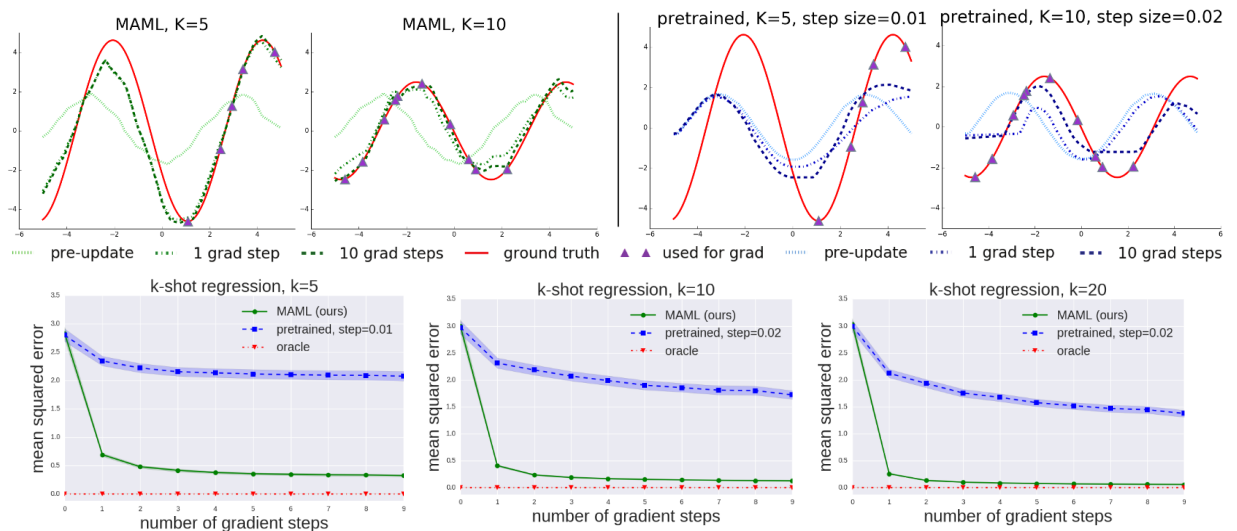


Figure 6. Quantitative sinusoid regression results showing test-time learning curves with varying numbers of K test-time samples. Each gradient step is computed using the same K examples. Note that MAML continues to improve with additional gradient steps without overfitting to the extremely small dataset during meta-testing, and achieves a loss that is substantially lower than the baseline fine-tuning approach.

- Even with only 5 datapoints the fitting curve is nice
- When all the points are in one half, the model can still infer the shape of the other half → model the periodic nature
- Quantitative results: improve with addition gradient steps
 - 但素, sin曲线这种简单任务好像也就是再多迭代一个循环的事情 :)
- 总之对Regression可以用很少的样本/循环finetune, 不会overfitting

5.2 Classification

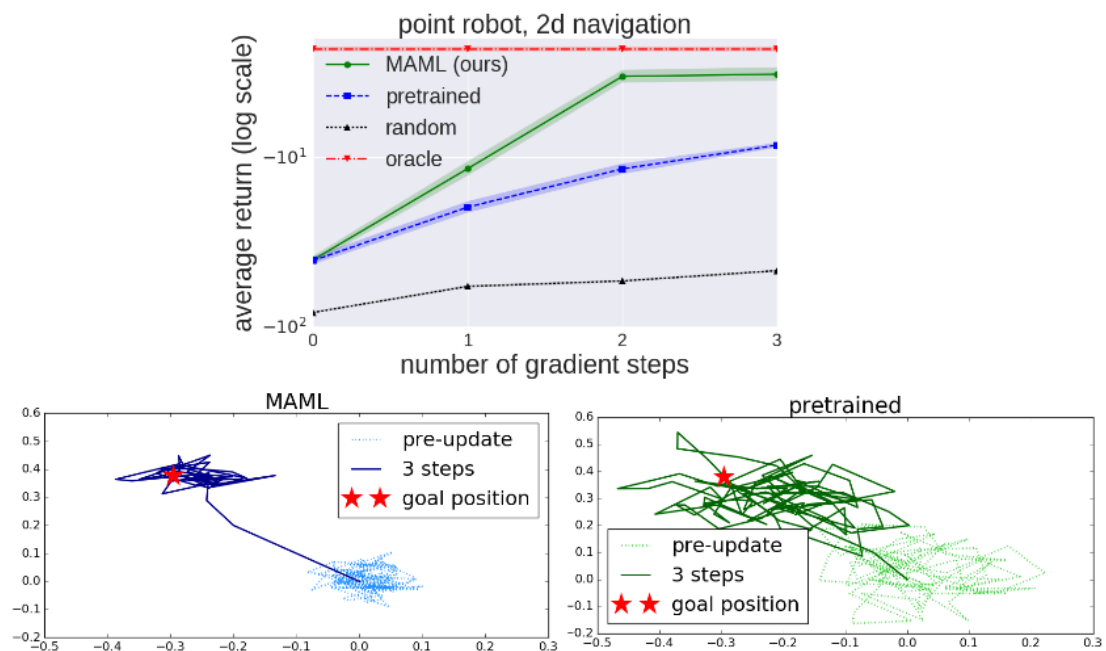
- Vinyals et al. 2016 Matching networks for one shot learning
- Tasks : Few-shot image recognition
- Datasets
 - Omniglot:
 - 20 instances of 1623 characters from 50 different alphabets
 - Downsampled to $28 * 28$
 - 1200 characters for training, remain for testing
 - Augmentation: degree rotations
 - MiniImagenet: 64 training classes, 12 validation classes, and 24 test classes
- Evaluation : N-way classification
 - Select N unseen classes

- Provide the model with K different instances of each of the N classes
 - Evaluate the model's ability to classify new instances within the N classes
- Model architecture
 - 4 modules, each module:
 - 3*3 conv, 64 filters (32 filters for MiniImagenet)
 - Batch normalization
 - ReLU nonlinearity
 - Strided convolutions (2*2 max-pooling for MiniImagenet)
 - A non-conv network for comparison: 256-128-64-64, BN, ReLU
 - Loss: cross entropy
- **Comparison result : see 4. related work**
- We also compare the performance between first order and second order derivatives
 - From $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{T_i \sim p(T)} L_{T_i}(f_{\theta_i})$ we can find that there is second order derivatives, which maybe computational expensive
 - Thus we compare it with first-order approximation
 - Compute the meta-gradient at the post-update parameter values θ'_i
 - Result:
 - **1st order is similar to 2nd order, less need to use 2nd derivatives**
 - Improvement of MAML comes from gradients of the objective at the post-update parameter values
 - Not 2nd derivative for differentiating through the gradient update
 - The use of ReLU make most of 2nd derivatives close to 0

5.3 Reinforcement Learning


- Dual et al. 2016 Benchmarking deep reinforcement learning for continuous control
- Model architecture
 - 2 hidden layers, 100 hidden nodes, ReLU
 - Vanilla PG + TRPO
 - Use finite differences to compute Hessian-vector products for TRPO: avoid computing third derivatives

- Comparison:
 - policy inited with MAML
 - policy inited with randomly weights
 - oracle policy
- 2D Navigation
 - Goal: move to goal positions in 2D
 - Observation: current 2D position
 - Action: velocity commands
 - Reward: negative squared distance
 - When to terminate: too close to the goal, or $H = 100$
 - Comparison: adaptation to a new task with up to 4 gradient updates, each with 40 samples



- MuJoCo Simulation:
 - Goal velocity experiment: the reward is the negative absolute value between the current velocity of the agent and a goal
 - Goal direction experiments: the reward is the magnitude of the velocity in either the forward or backward direction
 - Result : MAML learns a model that can quickly adapt its velocity and direction with even just a single gradient update, and continues to improve with more gradient steps

- 本文中经过MAML预训练的模型强于随机初始化, 不过在Parisotto et al 2016.的工作中, 也有预训练不如随机初始化的情况
- Note that random baseline curves for game velocity are removed (worse return)



6. Discussion and Future Work

- MAML can be treated as an initialization method to get pretrained model that is easy to fine-tune
- Benefits:
 - Simple, does not introduce any learned parameters for meta-learning
 - Can be applied to regression / classification / RL
 - Adaptation on new tasks with few shot / updates
- Future work:
 - Generalize meta-learning technique to apply to any problem and any model → 1709.04905 - One-Shot Visual Imitation Learning via Meta-Learning
 - Apply to multi-task

7. Code

- Here I try to understand `maml.py` in <https://github.com/cbfinn/maml>
- Note that this is only the code for SL, RL version is more complex
- How will it be used:
 - Suffix 'a': training data for inner loop
 - Suffix 'b': testing data for inner loop

```
1. model = MAML(dim_input, dim_output, test_num_updates =
   test_num_updates)
2. input_tensors = {'inputa': inputa, 'inputb': inputb, 'labela': labela,
   'labelb': labelb}
```

```
3. model.construct_model(input_tensors=..., prefix='...')
```

7.1 Initialize the model

- `update_lr` : learning rate α
- `meta_lr` : learning rate β
- `lossesa` : the training loss of inner loop \rightarrow for updating α
- `lossesb` : the testing loss of inner loop (the training loss of meta) \rightarrow for updating β

```
1. lossesa, outputas, accuraciesa = [], [], []
2. num_updates = max(self.test_num_updates, FLAGS.num_updates)
3. lossesb, outputbs, accuraciesb = [[]]*num_updates, [[]]*num_updates,
   [[]]*num_updates
```

7.2 Inner loop

- Helper function `task_metalearn(inp, reuse=True)`
 - Input : `inp = (self.inputa, self.inputb, self.labela, self.labelb)`
 - Output :

```
task_output = [task_outputa, task_outputbs, task_lossa, task_lossesb]
```

 - For classification, add another 2 elements

```
task_accuracya, task_accuraciesb
```
 - **Question: I can't find how `outputas, outputbs` will be used later**
 - Here we use another 2 helper functions:
 - `forward` : forward pass, get task outputs
 - `loss_func` : compute loss
- The first iteration and remaining are splitted:
 - `task_outputa` and `task_lossa` will only be computed in first iteration, in remaining iterations this will be computed as `loss` directly
 - Thus the final output of `task_outputa` and `task_lossa` will be computed in **first iteration**

```

1. # For first iteration:
2. task_outputa = self.forward(inputa, weights, reuse=reuse)
3. task_lossa = self.loss_func(task_outputa, labela)
4. # For remaining iterations:
5. loss = self.loss_func(self.forward(inputa, fast_weights, reuse=True),
    labela)

```

- Then we compute gradients for inner loop
 - Recall that we will make comparison between 2nd order derivation and 1st order

```

1. # For first iteration:
2. grads = tf.gradients(task_lossa, list(weights.values()))
3. # For remaining iterations:
4. grads = tf.gradients(loss, list(fast_weights.values()))
5. # if True --> only use 1st order
6. if FLAGS.stop_grad:
7.     grads = [tf.stop_gradient(grad) for grad in grads]
8. # Transfer to dict
9. gradients = dict(zip(weights.keys(), grads))

```

- Update the weights for sub-task: $\theta' \leftarrow \theta' - \alpha * grad$

```

1. fast_weights = dict(zip(weights.keys(), [weights[key] - self.update_lr*
    gradients[key] for key in weights.keys()]))

```

- Compute the test error for inner loop → the train error of meta process

```

1. output = self.forward(inputb, fast_weights, reuse=True)
2. task_outputbs.append(output)
3. task_lossbs.append(self.loss_func(output, labelb))

```

7.3 Meta update

- Note that meta update is only for **meta-training**
- Map `meta_tasklearn` to all data

```

1. result = tf.map_fn(task_metalearn, elems=(self.inputa, self.inputb, sel
    f.labela, self.labelb), dtype=out_dtype,
    parallel_iterations=FLAGS.meta_batch_size)

```

```

2.
3.     if self.classification:
4.         outputas, outputbs, lossesa, lossesb, accuraciesa, accuraciesb =
           result
5.     else:
6.         outputas, outputbs, lossesa, lossesb = result

```

- Compute **average loss** to update meta params

- `total_loss1` : the training loss of inner loop
- `total_losses2` : the testing loss of inner loop
- For classification we also need to compute accuracy

```

1.     self.total_loss1 = total_loss1 = tf.reduce_sum(lossesa) / tf.to_float(F
           LAGS.meta_batch_size)
2.     self.total_losses2 = total_losses2 = [tf.reduce_sum(lossesb[j]) / tf.to
           _float(FLAGS.meta_batch_size) for j in range(num_updates)]

```

- Meta update

- There are 2 kinds of update operations: `pretrain_op` and `metatraining_op`
- In `main.py`: `iter = pretrain_iter + metatraining_iter`, metatraining will only happen when it reaches `metatraining_iter`
- However in default settings these 2 iterations will not be together : one is 0 that there is only 1 kind of iteration

```

1.     # meta update for pretrain_op will be in all iterations
2.     self.pretrain_op =
           tf.train.AdamOptimizer(self.meta_lr).minimize(total_loss1)
3.
4.     # meta update for metatraining_op
5.     if FLAGS.metatraining_iterations > 0:
6.         optimizer = tf.train.AdamOptimizer(self.meta_lr)
7.         self.gvs = gvs = optimizer.compute_gradients(self.total_losses2[FLA
           GS.num_updates-1])
8.         if FLAGS.datasources == 'miniimagenet':
9.             gvs = [(tf.clip_by_value(grad, -10, 10), var) for grad, var in
           gvs]
10.         self.metatraining_op = optimizer.apply_gradients(gvs)

```

7.4 Meta testing

- `metaval_total_loss1` : training error of test task
- `metaval_total_losses2` : testing error of test task

```
1. self.metaval_total_loss1 = total_loss1 = tf.reduce_sum(lossesa) / tf.to_float(FLAGS.meta_batch_size)
2. self.metaval_total_losses2 = total_losses2 = [tf.reduce_sum(lossesb[j]) / tf.to_float(FLAGS.meta_batch_size) for j in range(num_updates)]
3. if self.classification:
4.     self.metaval_total_accuracy1 = total_accuracy1 = tf.reduce_sum(accuraciesa) / tf.to_float(FLAGS.meta_batch_size)
5.     self.metaval_total_accuracies2 = total_accuracies2 = [tf.reduce_sum(accuraciesb[j]) / tf.to_float(FLAGS.meta_batch_size) for j in range(num_updates)]
```