

1709.04905 - One-Shot Visual Imitation Learning via Meta-Learning

- **Yunqiu Xu**
 - Other reference:
 - Presentation: https://www.youtube.com/watch?v=_9Ny2ghjwuY
 - Code: <https://github.com/tianheyu927/mil>
-

1. Introduction

- Challenge: learning each skill from a scratch is infeasible
- One-Shot Visual Imitation Learning via Meta-Learning
 - Reuse past experience to train the "base model", then adapt it to new task with only a single demonstration
 - Visual: use raw visual inputs
 - Meta-Learning: MAML [C. Finn et.al. 2017](#)
- Prior work : take task identity / demonstration as the input into a contextual policy
- Our work : learn parameterized policy, then adapt into new tasks through a few gradient updates

2. Related work

- In this work, the state of environment is unknown → we feed raw sensory inputs to learn it
- Two challenges for learning from demonstrations then applying it to real world:
 - Compounding errors: not in this work
 - **The need of a large number of demonstrations for each task**
- Why don't use Inverse RL:
 - How does it work : recover reward function from demonstrations

- Pros: reduce demonstrations, better than behavioral cloning
- Cons:
 - Requires additional robot experience to optimize the reward [C. Finn et.al. 2016](#)
 - Hard to evaluate learned reward, especially for high-dim data (image)
 - Gan-based IRL (e.g. GAIL) is hard to train
- How do we reduce the demonstrations: **share data across tasks**
 - First, use a dataset of demonstrations of many other tasks for meta learning, in this way we can build a base model
 - Then we can adapt this base model to new task with only a few demonstrations
 - Meta-learning is similar to transfer learning to some extent, the different is not the transfer on dataset, but the transfer on task
 - Take a simple instance, if the robot is learned to pick apple, orange, pear ... then it can pick peach easily

3. Meta-Imitation Learning

- Goal : learn a policy that can quickly adapt to new tasks from a single demonstration of that task
- Each imitation task $T_i = \{\tau = \{o_1, a_1, \dots, o_T, a_T\} \sim \pi_i^*, L(a_{1:T}, \hat{a}_{1:T}), T\}$
 - τ : a demonstration generated by policy π_i^*
 - $L(a_1, \dots, a_T, \hat{a}_1, \dots, \hat{a}_T) \rightarrow R$: loss function to give feedback
 - **Note that this form is different from original MAML**

3.1 MAML

- Consider a policy π with parameter vector θ
- Sample a task T_i from $p(T)$
- Train this task with K samples (adapt π to T_i to get new parameter θ')
- Test this task, then treat the testing error as the training error of meta-process (Use $\theta'_1, \dots, \theta'_n$, to update θ)
- Meta objective:

$$\min_{\theta} \sum_{T_i \sim p(T)} L_{T_i}(f_{\theta'_i}) = \sum_{T_i \sim p(T)} L_{T_i}(f_{\theta - \alpha \nabla_{\theta} L_{T_i}(f_{\theta})}) \quad (1)$$

- Finally, you can adapt trained π to a new task with only a few data / gradient updates

Model-Agnostic Meta-Learning

Learn the weights Θ of a model such that gradient descent can make rapid progress on new tasks.

Algorithm 1 Model-Agnostic Meta-Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

```

1: randomly initialize  $\theta$ 
2: while not done do
3:   Sample batch of tasks  $\mathcal{T}_i \sim p(\mathcal{T})$ 
4:   for all  $\mathcal{T}_i$  do
5:     Evaluate  $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$  with respect to  $K$  examples
6:     Compute adapted parameters with gradient descent:  $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ 
7:   end for
8:   Update  $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$ 
9: end while

```




Figure 1. Diagram of our model-agnostic meta-learning algorithm (MAML), which optimizes for a representation θ that can quickly adapt to new tasks.

$$\mathcal{T}_i = \{\tau = \{\mathbf{o}_1, \mathbf{a}_1, \dots, \mathbf{o}_T, \mathbf{a}_T\} \sim \pi_i^*, \mathcal{L}(\mathbf{a}_{1:T}, \hat{\mathbf{a}}_{1:T}), T\}$$

Experts

3.2 Extend MAML to MIL

- \mathbf{o}_t is the observation at time t , i.e. an image, while \mathbf{a}_t is the action
- For demonstration trajectory τ , we use MSE to compute loss:

$$L_{T_i}(f_{\phi}) = \sum_{\tau_j \sim T_i} \sum_t \|f_{\phi}(\mathbf{o}_t^{(j)}) - \mathbf{a}_t^{(j)}\|_2^2 \quad (2)$$

- During meta-learning, we assume each task has at least 2 demonstrations, thus we can sample a set of tasks with two demonstrations per task
- Compute θ'_i with one demonstration \rightarrow inner loop of meta-learning
- Use another demonstration to "test" $\theta'_i \rightarrow$ update θ according to the gradient of meta-objective
- Meta-testing:
 - Sample a new task T and its one demonstration
 - This task can involve new goals or manipulating new, previously unseen

objects.

- Then we can adapt θ to this task

Algorithm 1 Meta-Imitation Learning with MAML

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

```
1: randomly initialize  $\theta$ 
2: while not done do
3:   Sample batch of tasks  $\mathcal{T}_i \sim p(\mathcal{T})$ 
4:   for all  $\mathcal{T}_i$  do
5:     Sample demonstration  $\tau = \{\mathbf{o}_1, \mathbf{a}_1, \dots, \mathbf{o}_T, \mathbf{a}_T\}$  from  $\mathcal{T}_i$ 
6:     Evaluate  $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$  using  $\tau$  and  $\mathcal{L}_{\mathcal{T}_i}$  in Equation (2)
7:     Compute adapted parameters with gradient descent:  $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ 
8:     Sample demonstration  $\tau'_i = \{\mathbf{o}'_1, \mathbf{a}'_1, \dots, \mathbf{o}'_T, \mathbf{a}'_T\}$  from  $\mathcal{T}_i$  for the meta-update
9:   end for
10:  Update  $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$  using each  $\tau'_i$  and  $\mathcal{L}_{\mathcal{T}_i}$  in Equation (2)
11: end while
12: return parameters  $\theta$  that can be quickly adapted to new tasks through imitation.
```

3.3 Two Head Structure

- Why use this: more flexibility during adapting
- The parameters of pre-update head are not used for post-update head in final
- Modification : parameters of final layers are not shared, forming two heads
 - Change loss function as:

$$L_{T_i}(f_{\phi}) = \sum_{\tau_j \sim T_i} \sum_t \|W y_t^{(j)} + b - a_t^{(j)}\|_2^2 \quad (3)$$

- $y_t^{(j)}$: post-synaptic activations of the last hidden layer
- W, b : weights and bias for last layer

- Then the meta-objective is about θ, W, b

$$\min_{\theta, W, b} \sum_{T_i \sim p(T)} L_{T_i}(f_{\theta'_i}) = \sum_{T_i \sim p(T)} L_{T_i}(f_{\theta - \alpha \nabla_{\theta} L_{T_i}(f_{\theta})}) \quad (4)$$

3.4 Learning to Imitate without Expert Actions

- Why use this : it is more practical to simply provide a video of the task being performed
- We just simplify this problem by simplify the loss function as

$$L_{T_i}(f_\phi) = \sum_{\tau_j \sim T_i} \sum_t \|W y_t^{(j)} + b\|_2^2 \quad (3)$$

- This can be a future question for more robust loss function

4. Network Architecture

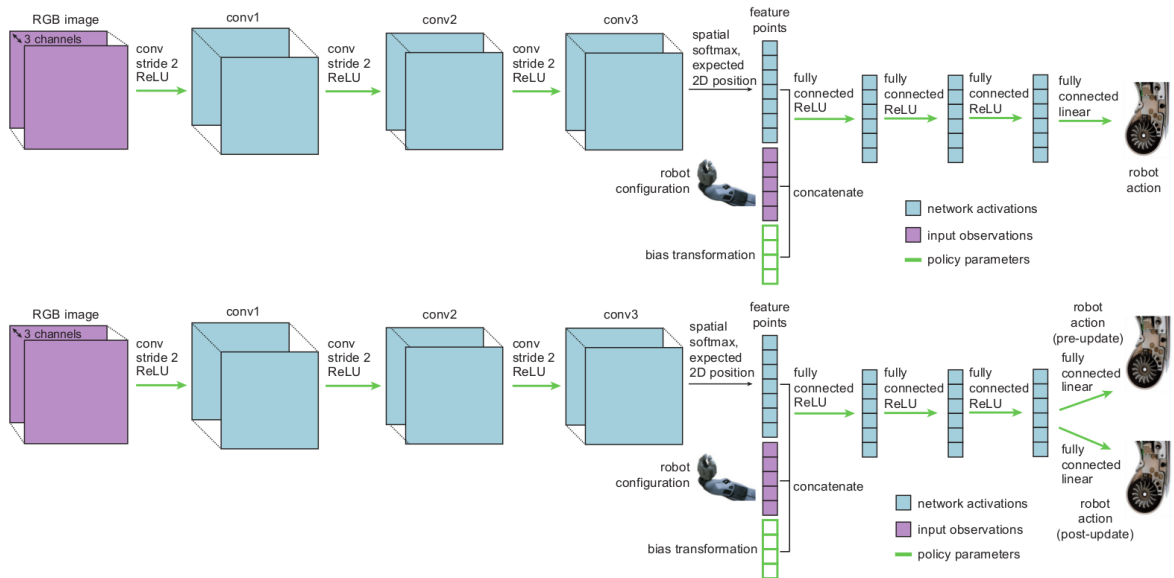


Figure 2: Diagrams of the policy architecture with a bias transformation (top and bottom) and two heads (bottom). The green arrows and boxes indicate weights that are part of the meta-learned policy parameters θ .

- **Layer normalization** after each layer
 - Data within a demonstration trajectory is highly correlated across time
 - Thus BN was not effective
- Bias transformation → improve the performance of meta-learning
 - Concatenate a vector of parameters to a hidden layer of post-synaptic activations
 - Thus vector is treated as same as other parameters during meta-learning and final testing


$$y = Wx + b \rightarrow y = W_1x + W_2z + b$$

- z is the parameter vector form bias transformation
- $W = [W_1, W_2]$
 - This modification **increases the representational power of the gradient**
 - Does not affect the representation power of the network itself

4. Experiment

- Questions:
 - Can a policy be learned to maps from image pixels to actions using a single demonstration of a task
 - How does our meta-imitation learning method compare to other one-shot imitation learning methods
 - Can we learn without expert actions
 - How well does our method scale to real world tasks
- Methods for comparison:
 - Our method
 - Random policy: output random actions from standard normal distribution
 - Contextual policy:
 - Input the final image of demonstration
 - Indicate goal and current image (observation)
 - Then output current action
 - LSTM:
 - Input demonstration and current observation
 - Output current action
 - LSTM + attention: only applicable to non-vision tasks
- Task 1 : simulated reaching
 - Goal: reach a target of a particular color
 - Both vision and non-vision versions are tested
 - meta-learning can handle raw-inputs
 - Our method can handle small dataset (demonstration) well
 - Bias transformation (bt) can perform more consistently across dataset sizes


Simulated reaching



Task:


- ☐ reaching a target of a particular color
- ☐ Policy must learn to localize the target using the demonstration and generalize to new positions
- ☐ Meta-training must learn to handle different colors
- ☐ 150 tasks x 10 different trials per task

Simulated reaching



- Task 2 : simulated pushing

Simulated pushing




Task:

- ❑ A push is considered as success if the center of the target object lands on the red target circle for at least 10 timestamps.
- ❑ Each task is defined as pushing a particular objects
- ❑ 74 tasks x 6 different trials per task

method		video+state +action	video +state	video
LSTM	1-shot	78.38%	37.61%	34.23%
contextual		n/a	58.11%	56.98%
MIL (ours)		85.81%	72.52%	66.44%
LSTM	5-shot	83.11%	39.64%	31.98%
contextual		n/a	64.64%	59.01%
MIL (ours)		88.75%	78.15%	70.50%

Table 1: One-shot and 5-shot simulating pushing success rate with varying demonstration information provided at test-time. MIL can more successfully learn from a demonstration without actions and without robot state and actions than LSTM and contextual policies.




- Task 3 : real-world placing
 - Experiment goal : place a held item into a target container, such as a cup, plate, or bowl, while ignoring two distractors

Real-World Placing



subset of training objects



test objects



subset of training objects



test objects


Task:

Evaluate how well a real robot (PR2) can learn to interact with new unknown objects from a single visual demonstration.

Success: the held object landed in or on the target container after the gripper is opened

method	test performance
LSTM	25%
contextual	25%
MIL	90%
MIL, video only	68.33%

Table 2: One-shot success rate of placing a held item into the correct container, with a real PR2 robot, using 29 held-out test objects. Meta-training used a dataset with ~ 100 objects. MIL, using video only receives the only video part of the demonstration and not the arm trajectory or actions.



5. Summary and Ongoing Work (On CoRL)

- Summary:
 - reuse prior experience when learning in new settings
 - learning-to-learn enables effective one-shot learning
- Ongoing work: one-shot imitation from human video → during demo, let human put the ball in cup