

1710.02298 - Rainbow: Combining Improvements in Deep Reinforcement Learning

RL AAAI

Author : Yunqiu Xu

0. Related resources

- http://www.sohu.com/a/197158891_610300
- <https://www.jiqizhixin.com/articles/2017-10-10-2>

1. Introduction

- Examines six extensions to the DQN algorithm as well as their combination
- Rainbow:
 - An agent that combines all the ingredients
 - Shows best performance on Atari 2600 benchmark, both in terms of data efficiency and final performance
 - Shows the contribution of each component to overall performance


2. Different DQNs

- Original DQN
 - NIPS 2013: [Playing Atari with Deep Reinforcement Learning](#)
 - Nature 2015: Human-level control through deep reinforcement learning
- Double DQN: [Deep Reinforcement Learning with Double Q-learning](#)
- Prioritized DQN: [Prioritized Experience Replay](#)

- Dueling DQN: [Dueling Network Architectures for Deep Reinforcement Learning](#)
- A3C: [Asynchronous Methods for Deep Reinforcement Learning](#)
- Distributional DQN: [A Distributional Perspective on Reinforcement Learning](#)
- Noisy DQN: [Noisy Networks for Exploration](#)

Network	What's new
DQN	CNN; experience replay; target network
DDQN	decoupling selection + evaluation of the bootstrap action → handle overestimation
Prioritization	Improves data efficiency
Dueling DQN	Advantage function
A3C	shifts the bias-variance tradeoff ; propagate newly observed rewards faster to earlier visited states
Distributed DQN	learns a categorical distribution of discounted returns, instead of estimating the mean
Noisy DQN	uses stochastic network layers for exploration

3. Comparison results:



Left: Rainbow and different DQNs. Right: when we remove some components