

1803.02999 - On First-Order Meta-Learning Algorithms

On First-Order Meta-Learning Algorithms

Alex Nichol and Joshua Achiam and John Schulman

OpenAI

{alex, jachiam, joschu}@openai.com

- Yunqiu Xu
-

1. Introduction

- Challenges:
 - Human can learn a new task quickly through prior knowledge
 - Meta learning can be used to achieve this, however MAML with two-level derivatives can be computational expensive
- Our contribution
 - Expand the work of First-order MAML, which is simpler to implement
 - Reptile: similar to FOMAML but doesn't need to split training-test for each task
 - Analyse both FOMAML and Reptile theritically to show that they optimize for within-task generalization
 - Experiment on Mini-ImageNet and Omniglot

2. Reptile

- 符号:
 - ϕ : 初始参数

- $\tilde{\phi}$: 梯度下降更新后参数
- $\tilde{\phi} = U_{\tau}^k(\phi)$: 使用任务 τ 更新 k 次参数 ϕ , 即进行 k 次梯度下降, 得到更新后的参数 $\tilde{\phi}$
- MAML:

Algorithm 1 Model-Agnostic Meta-Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

1: randomly initialize θ

2: **while** not done **do**

3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$

4: **for all** \mathcal{T}_i **do**

5: Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ with respect to K examples

6: Compute adapted parameters with gradient descent: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$

7: **end for**

8: Update $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$

9: **end while**

- MAML基于假设: 对于内循环(inner-loop)某个任务, 使用训练集 \mathbf{A} 进行训练, 并使用测试集 \mathbf{B} 进行测试, 其测试误差会被当作外循环(meta-loop)
- 可以看出MAML在内循环和外循环都需要计算梯度导数 \rightarrow 二阶导数

$$g_{\text{MAML}} = \frac{\partial}{\partial \phi} L_{\tau, B}(U_{\tau, A}(\phi)) \quad (3)$$

$$= U'_{\tau, A}(\phi) L'_{\tau, B}(\tilde{\phi}), \quad \text{where } \tilde{\phi} = U_{\tau, A}(\phi) \quad (4)$$

- FOMAML: 仅仅对外循环梯度求导, 不求取内循环梯度的导数, 将 $U'_{\tau, A}$ 看作常数
- Reptile:

Algorithm 1 Reptile (serial version)

Initialize ϕ , the vector of initial parameters
for iteration = 1, 2, ... **do**
 Sample task τ , corresponding to loss L_τ on weight vectors $\tilde{\phi}$
 Compute $\tilde{\phi} = U_\tau^k(\phi)$, denoting k steps of SGD or Adam
 Update $\phi \leftarrow \phi + \epsilon(\tilde{\phi} - \phi)$
end for

- 在内循环进行 k 次更新, 得到更新后的参数 $\tilde{\phi}$
- 然后用 $\tilde{\phi}$ 对 ϕ 进行更新

3. Experiment

- Few-shot classification:
 - Transductive setting: 通过batch normalization在test sample之间共享信息
 - Non-transductive setting: 仅仅在training sample上使用batch normalization, 使用 single test sample

Algorithm	1-shot 5-way	5-shot 5-way
MAML + Transduction	48.70 \pm 1.84%	63.11 \pm 0.92%
1 st -order MAML + Transduction	48.07 \pm 1.75%	63.15 \pm 0.91%
Reptile	47.07 \pm 0.26%	62.74 \pm 0.37%
Reptile + Transduction	49.97 \pm 0.32%	65.99 \pm 0.58%

Table 1: Results on Mini-ImageNet. Both MAML and 1st-order MAML results are from [4].

Algorithm	1-shot 5-way	5-shot 5-way	1-shot 20-way	5-shot 20-way
MAML + Transduction	98.7 \pm 0.4%	99.9 \pm 0.1%	95.8 \pm 0.3%	98.9 \pm 0.2%
1 st -order MAML + Transduction	98.3 \pm 0.5%	99.2 \pm 0.2%	89.4 \pm 0.5%	97.9 \pm 0.1%
Reptile	95.39 \pm 0.09%	98.90 \pm 0.10%	88.14 \pm 0.15%	96.65 \pm 0.33%
Reptile + Transduction	97.68 \pm 0.04%	99.48 \pm 0.06%	89.43 \pm 0.14%	97.12 \pm 0.32%

Table 2: Results on Omniglot. MAML results are from [4]. 1st-order MAML results were generated by the code for [4] with the same hyper-parameters as MAML.

- Different inner-loop and outer-loop gradient combinations
 - $\mathbf{g}_1, \mathbf{g}_2, \mathbf{g}_3, \mathbf{g}_4$: inner-loop gradients with different minibatches
 - $\mathbf{g}_1 + \mathbf{g}_2$: outer-loop update for two-step Reptile
 - \mathbf{g}_2 : outer-loop update for two-step FOMAML




Figure 3: Different inner-loop gradient combinations on 5-shot 5-way Omniglot.

- Overlap between inner-loop minibatches:
 - Aim: check whether small changes to optimization procedure can lead to large changes in performance
 - shared-tail (cycling): final inner-loop mini-batch comes from the same set as earlier inner-loop batches
 - separate-tail (**more correct**) : final mini-batch comes from a disjoint of data

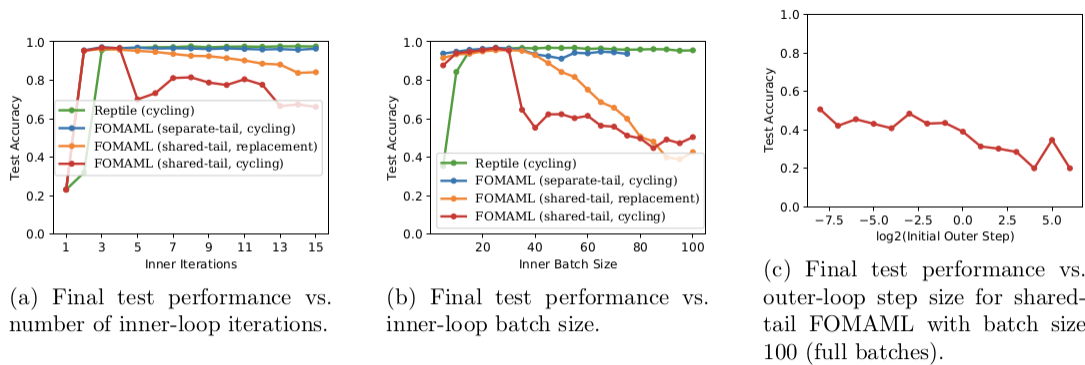


Figure 4: The results of hyper-parameter sweeps on 5-shot 5-way Omniglot.

- result: Reptile and FOMAML with cycling / separate-tail are not sensitive to inner-loop hyper-parameters

4. Summary

- Reptile:
 - 类似joint training和FOMAML
 - 将MAML简化为一阶
 - 和MAML差不多的性能, 但是效率更高
 - 和FOMAML相比, 对每个子任务不需要再split training-testing
- Future work:
 - Reptile在RL中的效果不大好, 作者这里提出原因可能为"joint training is a strong baseline", 未来需要对Reptile进行改进
 - 还有一些未来工作是关于few-shot classification的
- Further reading:
 - [Paper repro: Deep Metalearning using "MAML" and "Reptile"](#)
 - [Understanding Reptile: A Scalable Meta-learning Algorithm By OpenAI](#)
- Implementations:
 - [Pytorch implementation in "Paper repro: Deep Metalearning using 'MAML' and 'Reptile'"](#)
 - [TF implementation on new dataset](#)
 - [Another pytorch implementation](#)